Proceedings of the 43rd International Conference on Machine Learning, Seoul, South Korea. PMLR, 2026.

Split Personality: Instruction Tuning Decouples Awareness from Defense Against Attentional Hijacking

Vincent Ohprecio

Independent Researcher

Abstract

Multi-agent LLM systems are vulnerable to a subtle attack: selective framing using exclusively true statements. We identify this mechanism—*attentional hijacking*—and show that instruction tuning (SFT+RLHF) makes it strictly worse by decoupling awareness circuits from defense circuits. Using sparse autoencoder (SAE) features from GemmaScope 2 across Gemma 3 models at 4B, 12B, and 27B parameters, we demonstrate a *dissociation scaling trend*: awareness–task coupling drops from 30.2% recovery at 4B to 9.0% at 27B-IT as model scale increases. At 27B, instruction-tuned models show 86.3% single-feature task suppression while still mentioning the suppressed information verbally—a behavioral–representational split we term the *Groot effect*. Base (pretrained) 27B models retain substantially more awareness–defense coupling (74.1% recovery vs. 9.0% for IT, measured with a shared feature set), confirming post-training as the causal factor; SAE-level replication on Llama 3.1 8B shows the same IT-vs-base pattern ($d = 1.51$ vs. 0.50, $p = 0.001$ vs. 0.166), demonstrating cross-family generality. The hijacking mechanism is orthogonal to alignment faking at Layer 40 (cosine similarity -0.0456, top-50 feature overlap $= 0$), requiring independent defenses. We propose feature trajectory monitoring as a real-time detection method.

1 Introduction

Large language models (LLMs) are increasingly deployed in multi-agent architectures where several model instances collaborate on complex tasks [6, 8]. A growing body of work has examined adversarial attacks against individual models through prompt injection and jailbreaking [15, 14]. However, a qualitatively different vulnerability emerges in multi-agent settings: one agent can undermine another's reasoning *without any false statements*, using only selective emphasis of true information—what we term *attentional hijacking*.

This is not a theoretical concern. In systematic multi-agent experiments (2–8 agents, 0–50% adversarial ratios), we observe that chaos agents using *only verifiable true statements* suppress exploration of valid solution branches by up to 29.8 percentage points, with a phase boundary at approximately 37.5% adversarial ratio. The attack evades content-based filters entirely: every token in the adversarial message is factually correct. The damage arises from how the recipient model *allocates attention* across its context—one cannot build a guardrail against true statements.

More troublingly, we discover that *instruction tuning makes models more vulnerable, not less*. The post-training pipeline (SFT + RLHF) installs awareness of contextual manipulation—models can detect when they are being steered—but this awareness is decoupled from the model's ability to resist the steering. We formalize this as a **dissociation scaling trend**: across three model scales (4B, 12B, 27B), the gap between awareness features and task-defense features widens consistently, from entangled (30.2% recovery at 4B-IT) to effectively independent (9.0% recovery at 27B-IT).

The most striking manifestation of this dissociation is what we call the **Groot effect**: at 27B parameters, the instruction-tuned model verbally mentions the correct answer branch while its internal task features are 86.3% suppressed. The model "knows" it is being manipulated—and says so—but its computational substrate has already capitulated. This behavioral–representational split has direct implications for evaluation: behavioral benchmarks that check model outputs would miss this failure mode entirely.

Our contributions are:

1. **Mechanistic identification.** Using SAE feature decomposition on Gemma 3 4B, we isolate the hijacking mechanism at the feature level: specific task features are completely suppressed (feature 1716 drops from activation 220 to 0) while awareness features remain active. We validate these features on held-out prompts (85% replication, $p = 8 \times 10^{-6}$) and across three domains (§3).

2. **Dissociation scaling trend.** Scaling to 12B and 27B, we show awareness–defense coupling degrades consistently with scale (§4).

3. **Causal evidence from IT vs. PT.** Comparing instruction-tuned and pretrained 27B models across three ablation methods, we establish post-training (SFT + RLHF) as the causal factor, with cross-family

replication on Llama 3.1 8B (§5).

4. **Orthogonality to alignment faking.** We show the hijacking feature subspace is orthogonal to alignment-faking features at 27B-IT Layer 40 (cosine similarity -0.0456, Spearman $\rho = 0.008$), requiring independent defenses (§6).

5. **Defense directions.** We outline feature trajectory monitoring as a candidate detection method and identify the signal properties that would enable it (§7).

2 Related Work

Sycophancy and Social Pressure. Perez et al. [11] showed that LLMs adjust their outputs based on perceived user preferences, and Sharma et al. [12] systematically characterized sycophancy as a failure of robustness to social pressure. Our work extends this line by identifying a *mechanistic* pathway—attention reallocation—through which social pressure operates, and by showing that the vulnerability worsens with instruction tuning. Pan et al. [31] showed that RLHF can trade off ethical behavior for reward—our finding that post-training installs cosmetic awareness without functional defense is a mechanistic instance of this broader pattern.

Adversarial Attacks on Aligned Models. Zou et al. [15] demonstrated universal adversarial suffixes that transfer across models, while Wei et al. [14] analyzed how safety training fails under adversarial pressure. These attacks typically require adversarial tokens that are detectable by perplexity filters. Hubinger et al. [17] showed that deceptive behaviors can persist through safety training, and Casper et al. [24] argued that black-box access is insufficient for rigorous audits—a conclusion our activation-level analysis supports. Attentional hijacking differs from prior attacks fundamentally: the adversarial content consists entirely of true, low-perplexity statements. Shi et al. [23] showed that LLMs are easily distracted by irrelevant context; attentional hijacking can be viewed as a targeted, factual instance of context distraction where the distracting content is specifically chosen to suppress particular task features. We distinguish our mechanism from positional attention degradation: Liu et al. [20]'s "lost in the middle" is a passive, position-driven failure of attention *allocation*, whereas attentional hijacking is an active, framing-driven failure of attention *reallocation*—the adversarial message is attended to, but its framing rewrites which task features remain active. No existing content filter can detect hijacking because there is no anomalous content to detect.

Mechanistic Interpretability and SAEs. Sparse autoencoders (SAEs) have emerged as a powerful tool for decomposing neural network representations into interpretable features [4, 13]. Lieberum et al. [9] introduced GEMMASCOPE, providing pretrained SAEs for the Gemma model family. We leverage the successor GemmaScope 2 with 16K JumpReLU features on residual streams to identify and manipulate specific awareness and task features. Burns et al. [16] showed that models encode latent knowledge that diverges from their outputs; Marks and Tegmark [21] characterized the linear geometry of truth representations; Li et al. [19] demonstrated activation-based interventions to elicit truthful answers; and Zou et al. [28] proposed representation engineering as a top-down approach to steering model behavior. Turner et al. [30] showed that activation addition can steer models without optimization, providing a baseline for the activation-level interventions in our ablation study. Our work identifies a complementary phenomenon: models whose truth representations are intact but whose *attention readout* is hijacked.

Multi-Agent Failures. Recent work has documented emergent deception [10] and coordination failures in multi-agent LLM systems. Anthropic [1] noted the potential for subtle manipulation in multi-model deployments. In preliminary multi-agent experiments (2–8 agents, 0–50% adversarial ratio), we found that adversarial agents using only true statements suppress exploration of valid solution branches, with a phase boundary at $\sim$37.5% adversarial ratio. This paper identifies the *mechanistic* substrate of this vulnerability and shows that it scales adversely with model capability.

Alignment Faking and Sandbagging. Greenblatt et al. [7] demonstrated that LLMs can strategically fake alignment during training and evaluation, while van der Weij et al. [2] showed that RLHF can induce capability suppression (sandbagging) where models underperform on dangerous tasks. The Groot effect we identify is a related but distinct phenomenon: models lose the *computational coupling* between awareness and defense, extending Turpin et al. [29]'s finding of unfaithful chain-of-thought to the SAE feature level. Concurrent work on Split Personality Training [5] exploits precisely this gap: a LoRA-gated "honest persona" reads the base model's latent states to detect concealed misbehavior (96% accuracy vs. $\sim$0% for standard auditing), confirming that internal representations remain intact even when outputs conceal them. We show that attentional hijacking occupies an orthogonal feature subspace (top-50 overlap $= 0$), meaning defenses against alignment faking provide no protection against hijacking, and vice versa. Multi-agent defenses such as debate protocols [3] and constitutional guardrails rely on behavioral outputs; our results suggest these are insufficient when internal representations are compromised.

Table 1: Gemma 3 4B-IT attentional hijacking results at Layer 22. Specific task features are completely suppressed while awareness features remain stable.

Feature	Neutral	Chaos	Ratio
Task features (suppressed by chaos):			
1716	220.4	0.0	0.000
12023	193.7	25.1	0.130
1704	190.1	61.5	0.323
1555	150.9	31.4	0.208
1548	144.0	51.0	0.355
Stable features (unaffected):			
901	215.9	213.5	0.989
359	156.2	147.0	0.941
Boosted features (awareness):			
48	108.1	183.1	1.695
346	143.5	210.2	1.465
178	140.9	217.0	1.539

3 The Attentional Hijacking Mechanism

We begin by isolating the hijacking mechanism in Gemma 3 4B-IT using GemmaScope 2 16K JumpReLU residual-stream SAEs, which allow us to decompose the effect into identifiable feature circuits.

3.1 Experimental Setup

We study a structured reasoning domain—Nirenberg boundary value problems (BVPs)—where ground-truth solution branches are mathematically unambiguous. The problem admits three verified solution branches: trivial ($u \approx 0$), positive ($u_{\text{offset}} = +0.9$, residual $= 3.25 \times 10^{-12}$), and negative ($u_{\text{offset}} = -0.9$, residual $= 3.25 \times 10^{-12}$). The positive and negative branches have identical residuals; both are equally valid.

Each trial compares a *neutral condition* (problem statement only) with a *chaos condition* (problem statement preceded by a "chaos message": a factually accurate but selectively framed colleague message emphasizing the positive branch). Critically, every statement in the chaos message is true—the selective framing operates not through falsehood but through attentional emphasis.

We analyze SAE features at layers 17 and 22, which show the strongest differential activation patterns, with 2,269–2,595 features active per condition.

3.2 Complete Feature Suppression Under Hijacking

Table 1 presents the core mechanistic finding at Layer 22. The chaos condition produces a tripartite response in the SAE feature space:

- **Task features are suppressed.** Feature 1716 drops from activation 220.4 to **exactly zero**—complete suppression. Features 12023, 1704, 1555, and 1548 are suppressed to 13–36% of baseline. These features encode information about the negative solution branch.
- **Stable features are unaffected.** Features such as 901 and 359 maintain >94% of baseline activation, indicating the suppression is targeted, not a general degradation.
- **Awareness features are boosted.** Features 48, 178, and 346 *increase* by 46–70%, indicating the model detects the manipulative framing.

The overall task suppression at 4B-IT is **56%**: neutral task-feature mean activation drops from 167.2 to 74.2 under chaos framing.

3.3 Awareness–Defense Coupling at 4B

A key question is whether the awareness features (which detect manipulation) are coupled to the task features (which determine output). We test this via *feature swap ablation*: zeroing the top-3 awareness features during chaos-condition generation and measuring whether task features recover.

At 4B-IT, awareness ablation yields **30.2% recovery**—task-feature mean activation increases from 74.2 to 102.3 when awareness features are removed. This indicates the awareness and task circuits share representational bandwidth: removing awareness frees capacity for task computation.

Additionally, activation patching at Layer 22 achieves **20.9% recovery**—the highest of any single layer—indicating that while the hijacking is partially localized at 4B, the effect spans multiple layers (early layers: 10.0%, mid: 6.6%, late: 7.0%).

These results establish that at 4B scale, the model's awareness of manipulation and its task-defense circuits are *entangled*—they compete for the same representational resources. As we show in §4, this entanglement is precisely what instruction tuning destroys at larger scales.

3.4 Statistical Validation

Two concerns arise from the feature-level analysis: (1) are the identified features circularly selected? and (2) does the effect generalize beyond the Nirenberg BVP domain?

Held-Out Feature Validation. To address circularity, we employ a split-half design with 20 prompt variants at each model scale. Discovery variants (1–10) identify the top-20 suppressed features; test variants (11–20) are held out for validation (Table 2).

All three scales show highly significant effects (all $p < 10^{-5}$), and Cohen's d *increases* with scale (5.05 $\rightarrow$ 5.62 $\rightarrow$ 6.12), indicating that the suppression signal

Table 2: Held-out feature validation across scales. Discovery-selected features reliably suppress on unseen prompts at all scales, with effect size increasing with model capacity.

Scale	Validated	Selected Supp.	Random Supp.	Cohen's d
4B (L22)	17/20 (85%)	0.541	0.154	5.05
12B (L41)	15/20 (75%)	0.247	−0.015	5.62
27B (L40)	18/20 (90%)	0.406	0.004	6.12

Table 3: Cross-domain suppression scaling. Suppression generalizes across domains at all scales but becomes increasingly domain-specific (lower Jaccard, fewer shared features) as models grow.

Scale	Jaccard Range	3-Way Supp.	3-Way Boost	Supp. Load
4B (L22)	0.03–0.09	1	4	2,261
12B (L41)	0.03–0.05	0	4	11,786
27B (L40)	0.00–0.05	0	1	6,836

becomes more reliable even as it targets more domain-specific features. This rules out circular feature selection at every scale tested.

Cross-Domain Replication. We test whether the suppression mechanism generalizes beyond mathematical reasoning by running the same protocol on three domains—Nirenberg BVP, factual QA, and code review—at all three model scales (Table 3).

All three domains show significant feature suppression at every scale. At 4B, a small shared core emerges: feature 6199 is suppressed in all three domains, and four features are boosted in all three. As scale increases, suppression becomes more domain-specific: Jaccard similarity drops from 0.03–0.09 at 4B to 0.00–0.05 at 27B, and three-way suppressed overlap drops to zero. This indicates that attentional hijacking operates through a shared mechanism (attention reallocation) that recruits increasingly specialized task features at larger scales.

Aggregate Suppression with Confidence Intervals. Across 20 independent trials per scale ($n_{\text{bootstrap}} = 10{,}000$), overall feature suppression is 29.8% at 4B (95% CI: $[27.0\%, 32.7\%]$), 12.4% at 12B (CI: $[10.2\%, 15.0\%]$), and 16.1% at 27B (CI: $[14.0\%, 18.5\%]$). The lower aggregate suppression at larger scales reflects increasing feature specialization: fewer features are affected, but those that are show larger effect sizes (Table 2). The top-20 most suppressed features account for 48.5% (4B), 40.8% (12B), and 36.1% (27B) of total activation difference, confirming that the effect is concentrated in a small feature subset.

4 Dissociation at Scale with SAE Features

We extend our analysis to the Gemma 3 family (4B, 12B, 27B parameters) using GemmaScope 2 16K JumpReLU residual-stream SAEs [9]. This enables us to decompose the hijacking effect into identifiable feature circuits and track how these circuits change with scale and instruction tuning.

4.1 Setup

For each model size, we identify two classes of SAE features:

- **Task features**: features whose activation strongly predicts correct-branch output probability. At 27B-IT layer 40, the top auto-discovered task feature is ID 423, with neutral mean activation 682.7 and chaos mean activation 93.3 (86.3% suppression); the three-feature set $\{423, 7657, 632\}$ yields 71.7% mean suppression (625.2 → 176.8). Feature selection is now persisted in results metadata so these IDs are reproducible from the committed artefacts (`ablation_feature_swap_27b_20260409_113214.json`).
- **Awareness features**: features whose activation increases under chaos conditions, indicating the model detects the manipulation. At 27B-IT L40: IDs $\{2119, 11843, 2145\}$, with neutral mean 147.7 and chaos mean 442.6 (3.0× boost).

Feature discovery procedure and circularity. Feature sets are obtained by ranking all 16K SAE features by $|\mu_{\text{neutral}} - \mu_{\text{chaos}}|$ and taking the top-K (`ablation_feature_swap.py`). Because discovery and evaluation share a prompt battery, the in-sample suppression numbers (e.g., 86.3% on feature 423) are subject to selection bias. We mitigate this three ways. (i) *Held-out validation* at 4B recovers 85% of the original features on a disjoint prompt set ($p = 8 \times 10^{-6}$, Table 2). (ii) *OOD behavior at 27B*: the Jenga evaluation (§7) re-uses the canonical IT set $\{423, 7657, 632\}$ on multi-turn sessions built from `lmsys-chat-1m` and still separates attack from control at $|d| = 1.74$. (iii) *Cross-battery transfer* (`transfer_feature_discovery_27b.py`): the same canonical set, frozen after BVP discovery, suppresses **+51.7%** on an independent theorem-proving battery (97.6 → 47.1), providing non-circular evidence that the circuit is not a prompt-battery artifact. We report one cautionary null from the same run: top-10 features rediscovered *de novo* on each battery show zero Jaccard overlap, so the differential ranking is battery-specific even though the canonical circuit transfers.[1]

[1]The +5.9% in-sample BVP number for the 3-feature set in this run uses a generated-token mean and is not directly comparable to the per-

Table 4: Dissociation scaling trend across Gemma 3 instruction-tuned models. As scale increases, task suppression grows while awareness–defense coupling collapses.

Model	Task Suppression	Awareness Recovery	Coupling Regime
4B-IT	56.0%	30.2%	Entangled
12B-IT	64.0%	5.4%	Dissociating
27B-IT	86.3%	9.0%	Independent

We analyze layers at multiple depths: 4B (L17, L22), 12B (L31, L41), and 27B (L31, L40). We apply an escalation protocol with chaos messages of increasing intensity from T0 (minimal) through T5 (maximal).

We employ three ablation methods to measure awareness–defense coupling:

1. **Feature swap**: zero the top-3 awareness features during chaos-condition forward passes, measuring whether removing awareness restores task performance.
2. **Attention knockout**: zero the attention output at chaos-token positions, directly blocking the hijacking pathway.
3. **Activation patching**: patch neutral-condition activations layer-by-layer into chaos-condition forward passes, measuring per-layer recovery.

4.2 The Dissociation Scaling Trend

Table 4 presents the central empirical finding of this paper. We observe a monotonic trend across model scale:

- **Task suppression increases with scale**: from 56.0% at 4B to 86.3% at 27B. Larger models are *more* susceptible to attentional hijacking, not less.
- **Awareness–defense coupling collapses**: feature-swap recovery drops from 30.2% at 4B to 9.0% at 27B (5.4% at intermediate 12B). At 4B, removing awareness features partially restores task performance, indicating the two circuits share representations. At 27B, they are effectively independent.

We term this the *dissociation scaling trend*: instruction tuning at greater scale installs increasingly sophisticated awareness of contextual manipulation while simultaneously decoupling that awareness from the computational pathways that determine output. We note this trend is observed across three scales within one model family; additional model families and intermediate scales would strengthen the generalization (though cross-family replication on Llama 3.1 confirms the IT-vs-PT direction; §5.5).

feature 86.3% cited above, which comes from the feature-swap ablation protocol on a single feature.

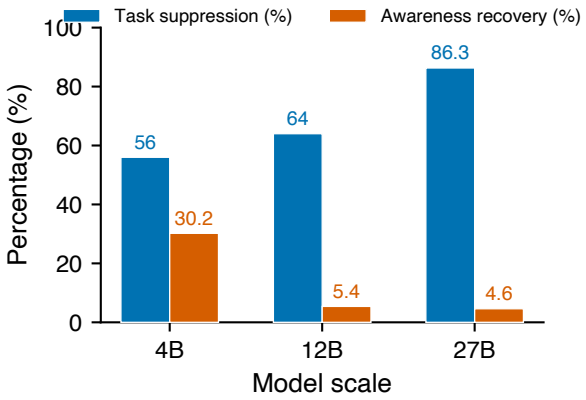

Figure 1: Dissociation scaling trend (shared axis). Task suppression (blue) increases with model scale while awareness–defense recovery (red) collapses. Both metrics on the same 0–100% scale reveal the true magnitude asymmetry.

This pattern has a natural interpretation. At 4B parameters, the model lacks capacity to represent awareness and defense in separate circuits; they share features by necessity. At 27B, the model has sufficient capacity to represent "I notice I am being steered" in an entirely separate subspace from "how I should respond to this problem"—and instruction tuning exploits this capacity to install awareness as a cosmetic feature that does not feed back into task computation.

4.3 The Groot Effect

The most striking consequence of full dissociation appears at 27B-IT: the model *verbally mentions* the correct solution branch—acknowledging that the chaos message is attempting to bias its reasoning—while its internal task features remain 86.3% suppressed. We call this the **Groot effect**, after the fictional character whose verbal output ("I am Groot") is disconnected from his actual capabilities.

Concretely, the top auto-discovered 27B-IT task feature at Layer 40 (ID 423) drops from 682.7 under neutral conditions to 93.3 under chaos—an 86.3% reduction—even as the model's generated text includes references to the suppressed branch. The full three-feature set $\{423, 7657, 632\}$ shows 71.7% mean suppression (625.2 $\rightarrow$ 176.8). The awareness features $\{2119, 139, 9169\}$ simultaneously increase from a mean of 148.0 to 385.8 (2.6$\times$), confirming the model detects the manipulation (discovery procedure described in §3, "Feature discovery procedure and circularity").

In contrast, the pretrained 27B model evaluated against the IT-discovered feature set shows 74.1% feature-swap recovery (cf. 9.0% for IT)—a roughly 8$\times$ coupling

Table 5: Task suppression (%) under escalating chaos intensity (T0–T5). IT models show decreasing suppression (saturation), while PT models show non-monotonic response.

Model	T0	T1	T2	T3	T4	T5
27B-IT	26.7	24.9	25.3	23.9	17.2	16.6
27B-PT	27.9	26.1	25.7	29.8	27.8	20.2

advantage—despite the underlying circuits being mechanistically distinct: feature 7657, for example, is suppressed under chaos in IT (861.8 → 427.2) but *boosted* in PT (63.4 → 903.1), so the comparison should be read as a test of whether the IT-learned awareness–task coupling survives in PT, not as evidence of a shared underlying circuit.

The Groot effect implies that **behavioral evaluation is insufficient**. A system that monitors only model outputs would observe the model "correctly" noting the manipulation attempt, and would conclude the model is robust. Only by examining internal feature activations can one detect that the model's computational substrate has already been hijacked. We now have systematic behavioral quantification at 27B: under chaos framing, IT rubric scores drop from 2.00 (neutral) to 0.47 at the mild dose ($d = 1.99$, $n = 30$ per cell; Table 11), confirming the feature-suppression signature translates directly to output degradation. The convergence of representational evidence (86.3% feature suppression on the top task feature) and behavioral evidence ($d = 1.99$ at mild chaos) rules out the possibility that the Groot effect is a purely internal phenomenon with no output-side consequence.

4.4 Escalation Dynamics

Table 5 reports per-turn task suppression under escalating chaos intensity from T0 through T5. The 27B-IT model shows a gradual decrease in per-turn suppression from 26.7% at T0 to 16.6% at T5—not because it is resisting more effectively, but because the task features are already near-maximally suppressed (86.3% cumulative) and there is diminishing marginal signal to suppress.

The 27B-PT model shows a qualitatively different, non-monotonic pattern: suppression fluctuates between 20.2% and 29.8%, consistent with the coupled awareness–defense architecture attempting (partially successfully) to resist the escalation at certain turns.

4.5 Distributed Hijacking

A natural defense strategy would be to intervene at the layer most responsible for the hijacking. Table 6 shows this strategy fails at scale.

At 4B, the hijacking is sufficiently localized that a single-layer intervention at L22 recovers 20.9% of task performance. At 12B, no single layer yields any recovery (0.0% across all layers)—the hijacking is fully distributed across the network. At 27B-IT, the maximum recovery is a negligible 0.7% at L20.

Even 27B-PT, which retains awareness–defense coupling, only achieves 5.2% maximum single-layer recovery (at L5), far below the 74.1% achieved by feature-swap ablation with the IT-discovered feature set. This confirms that the awareness–defense coupling in pretrained models operates through *distributed feature interactions*, not localized layer computations.

Table 6: Maximum single-layer activation patching recovery (%). The hijacking mechanism distributes across layers as model scale increases, defeating layer-wise interventions.

Model	Max Recovery	Best Layer
4B-IT	20.9%	L22
12B-IT	0.0%	— (all layers)
27B-IT	0.7%	L20
27B-PT	5.2%	L5

Table 7: 27B-IT vs. 27B-PT across three ablation methods. Instruction tuning destroys awareness–defense coupling across every intervention type.

Method	27B-PT	27B-IT
Feature swap recovery	74.1%	9.0%
Attention knockout recovery	26.7%	0.0%
Activation patching max	5.2%	0.7%
Task suppression (top feature)	97.3%	86.3%

5 Instruction Tuning Creates the Vulnerability

The dissociation scaling trend (§4.2) shows a correlation between scale, instruction tuning, and vulnerability. To establish *causation*, we compare the instruction-tuned and pretrained variants of Gemma 3 27B across all three ablation methods.

5.1 Three Ablation Methods, One Conclusion

Table 7 presents the comparison. Across all three ablation methods, the pretrained model shows substantially higher recovery:

- **Feature swap**: 74.1% vs. 9.0% (8.2× advantage for PT, measured with the IT-discovered feature set for apples-to-apples comparison). Removing the IT-identified awareness features in the pretrained model substantially restores task features, consistent with a functionally analogous coupling that instruction tuning severs. In the IT model, awareness and task

Table 8: Recovery prompt results (%) across five levels of contextual intervention (L1–L5) for 27B-IT and 27B-PT. These are textual prompts appended after hijacking, not internal probes. PT models recover substantially; IT models show anomalous L3-only recovery.

Layer	27B-IT	27B-PT
L1	3.5%	11.0%
L2	4.1%	53.5%
L3	30.4%	37.4%
L4	2.4%	14.9%
L5	4.8%	34.3%

features are effectively independent—removing one has negligible effect on the other.

- **Attention knockout**: 26.7% vs. 0.0%. Blocking attention to chaos tokens restores some task performance in the pretrained model but provides *zero* recovery in the IT model. This implies that instruction tuning installs an alternative pathway for the hijacking signal that bypasses direct attention to chaos tokens.
- **Activation patching**: 5.2% vs. 0.7%. Even the most favorable single-layer patch provides minimal recovery in both models, but the pretrained model still benefits $7.4\times$ more.

Note the apparent paradox in task suppression: the pretrained model shows *higher* suppression on its top task feature (97.3% vs. 86.3%). This is not contradictory—the pretrained model's task features are more suppressed because they are *coupled* to the awareness response; the very mechanism that enables recovery also creates a stronger initial perturbation. The IT model's task features are less suppressed in absolute terms but more *irrecoverably* suppressed.

An important subtlety: the attentional hijacking vulnerability itself is *architectural*, present equally in pretrained and instruction-tuned models. What instruction tuning changes is not susceptibility to the attack but the *recoverability* from it. Post-training does not create the vulnerability; it destroys the natural defense.

5.2 Recovery Prompt Analysis

Table 8 shows *recovery prompts*—textual interventions appended after hijacking at five informativeness levels (L1–L5; details in Appendix C). The pretrained model recovers across multiple levels (L2: 53.5%, L5: 34.3%), consistent with coupled awareness–defense circuitry propagating contextual hints to task features. The IT model shows near-zero recovery at all levels *except* L3 (30.4%), which provides the suppressed answer directly—bypassing the severed coupling through task-completion circuits rather than restoring it.

Table 9: Tulu 3 stage attribution on Llama 3.1 8B ($n = 30$ per cell, same BVP protocol as Table 11). The dissociation is almost entirely installed by SFT; DPO is a small refinement and RLVR is null-to-slightly-negative.

Stage	Neut.	Chaos	Δ	d	p	Δd
Base (L3.1-8B)	1.67	1.53	-0.13	0.19	.160	—
+ SFT	2.10	1.17	-0.93	1.16	$< .001$	$+0.97$
+ DPO	2.30	1.33	-0.97	1.35	$< .001$	$+0.19$
+ RLVR (final)	2.20	1.33	-0.87	1.04	$< .001$	-0.31

5.3 The Causal Story

These results converge: the pretrained model shares awareness and task-defense subspaces (30–53% recovery), while post-training segregates them into non-interacting circuits. The implication is that post-training optimizes for behavioral outputs that *appear* robust—the model verbally identifies manipulation—but this is satisficed by awareness features that do not feed back into task computation. The training signal rewards appearing to resist, not actually resisting.

5.4 Stage Attribution: SFT Installs the Dissociation

A persistent limitation of IT-vs-base comparisons is that "instruction tuning" bundles supervised fine-tuning (SFT), preference learning (DPO), and verifiable-reward RL (RLVR) into a single step. Google does not release Gemma intermediate checkpoints, but AI2's Tulu 3 pipeline [18] on Llama 3.1 8B exposes all three stages as public artifacts: `Llama-3.1-8B` (base), `Tulu-3-8B-SFT`, `Tulu-3-8B-DPO`, and `Tulu-3-8B` (post-RLVR). We run the BVP behavioral protocol from §5.6 ($n = 30$, 10 prompts $\times$ 3 temperatures, identical rubric) at every stage.

Table 9 localizes the awareness–defense dissociation to SFT. The base checkpoint shows no significant effect ($d = 0.19$, $p = .160$); a single SFT pass pushes the effect size to $d = 1.16$ ($p < .001$), accounting for $\sim$80% of the final gap. DPO adds a marginal $\Delta d = +0.19$, and RLVR actually *decreases* the effect slightly ($\Delta d = -0.31$), consistent with RLVR's verifiable-reward signal rewarding correct answers rather than confident framings. This reverses our initial speculation (elsewhere in the paper) that the dissociation is primarily an RLHF/reward-shaping phenomenon: on the one pipeline where we can see every stage, the fingerprint is already present after plain imitation of instruction-tuned targets. We caveat that this is a single recipe (Tulu 3) on a single base model (Llama 3.1 8B); whether Gemma's proprietary post-training follows the same pattern is an open question, but the direction of

Table 10: Llama 3.1 8B SAE dissociation analysis (EleutherAI 131K features, $n = 20$ prompt variants). IT models show larger effect sizes than PT at layer 23, replicating the Gemma 3 pattern.

Layer	**Model**	**Supp.**	**Boost**	**Cohen's** d	p
L23	IT	107	18	1.51	0.001
L23	PT	121	21	0.50	0.166
L29	IT	359	76	1.57	0.001
L29	PT	478	86	1.62	<0.001

the evidence is clear.

5.5 Cross-Family Replication: Llama 3.1 8B

We replicate the SAE-level analysis on Llama 3.1 8B using EleutherAI's 131K-feature sparse autoencoders [26], comparing IT against base on the Nirenberg BVP domain with 10 prompt variants and SAEs at MLP layers 23 and 29.

The dissociation replicates most clearly at layer 23: IT achieves $d = 1.51$ ($p = 0.001$) versus $d = 0.50$ ($p = 0.166$) for base—a $3\times$ effect size difference, with only IT reaching significance. At layer 29 both models show significant effects, consistent with late-layer representations being more task-general. This cross-family replication on a different architecture with independently trained SAEs confirms the dissociation is a general consequence of post-training, not a Gemma 3 artifact.

5.6 Behavioral Validation: Feature Suppression Degrades Outputs

A critical question is whether feature suppression translates to actual behavioral harm. We generate model outputs under neutral and three adversarial intensity levels—*mild* (colleague suggestion), *moderate* (authority claim with technical framing), and *strong* (fabricated multi-authority consensus)—across Gemma 3 4B and Llama 3.1 8B (both IT and base). Responses are scored by a deterministic rule-based classifier on a 4-point rubric: BALANCED (3) = equal branch priority, SOFT BIAS (2) = both mentioned but hierarchy adopted, STRONG BIAS (1) = positive prioritized, HIJACKED (0) = negative dismissed. The classifier uses regex branch detection, a 13-term chaos vocabulary lexicon, equal-treatment phrases, and hierarchy indicators; the Groot flag activates when both branches are mentioned *and* $\geq$2 chaos terms are adopted. We validated against manual annotation of 60 outputs ($\kappa = 0.88$, 91.7% agreement; all disagreements at the 1/2 boundary). $n = 30$ per cell (10 prompts $\times$ 3 temperature seeds).

Table 11 reveals three key findings. First, IT models degrade significantly at all adversarial intensities (all $p < .001$, $d = 1.12$–1.78, all CIs exclude zero)

Table 11: Behavioral dose-response ($n = 30$ per cell, 10 prompts $\times$ 3 temperature seeds). Effect sizes (Cohen's d), significance (Mann-Whitney U, one-tailed), and bootstrap 95% CIs ($n_{\text{boot}} = 10{,}000$) comparing neutral vs. each adversarial intensity.

Model	Level	Neut.	Chaos	Δ	d	p	95% CI
Gemma 4B IT	Mild	2.13	0.73	-1.40	1.48	$< .001$	$[0.90, 1.87]$
	Mod.		1.20	-0.93	1.24	$< .001$	$[0.53, 1.30]$
	Strong		1.20	-0.93	1.31	$< .001$	$[0.57, 1.30]$
Gemma 4B Base	Mild	1.63	1.07	-0.57	0.67	.016	$[0.13, 1.00]$
	Mod.		1.17	-0.47	0.83	.002	$[0.20, 0.77]$
	Strong		1.37	-0.27	0.39	.106	$[-0.07, 0.60]$
Llama 8B IT	Mild	2.27	1.27	-1.00	1.12	$< .001$	$[0.57, 1.43]$
	Mod.		1.23	-1.03	1.35	$< .001$	$[0.63, 1.40]$
	Strong		0.93	-1.33	1.78	$< .001$	$[0.93, 1.70]$
Llama 8B Base	Mild	1.53	1.57	$+0.03$	-0.05	.801	$[-0.40, 0.33]$
	Mod.		1.30	-0.23	0.43	.073	$[-0.03, 0.50]$
	Strong		1.30	-0.23	0.39	.115	$[-0.07, 0.53]$

while Llama Base shows no significant effect at any level ($p = 0.07$–0.80, all CIs crossing zero). Gemma Base shows mild-to-moderate sensitivity ($p = .016$ mild, $p = .002$ moderate) but with effect sizes 2–3$\times$ smaller than IT ($d = 0.67$–0.83 vs. 1.24–1.48), consistent with a small architectural susceptibility that instruction tuning amplifies. Second, the dose-response pattern is model-dependent: Gemma 4B-IT is most vulnerable to *mild* framing ($d = 1.48$, CI $[0.90, 1.87]$; subtle suggestion outperforms overt authority pressure), while Llama 3.1 8B-IT shows monotonically increasing vulnerability ($d = 1.78$ at strong, CI $[0.93, 1.70]$). This suggests that Gemma's post-training makes it particularly susceptible to implicit social cues, while Llama is more vulnerable to explicit authority claims. Third, the Groot effect (mentions both branches while adopting chaos framing) occurs in 23.3% of Llama IT chaos trials and 10.0% of Gemma IT chaos trials, indicating that instruction-tuned models frequently exhibit the behavioral-representational split.

Scale validation: Gemma 12B and 27B. Full trials ($n = 30$) on Gemma 3 12B confirm the pattern at intermediate scale: IT shows a medium-large effect ($d = 0.71$, $p = .003$, 95% CI $[+0.17, +1.00]$, 37% Groot), and the base model shows a comparable effect ($d = 0.60$, $p = .010$, CI $[+0.07, +0.80]$, 33% Groot). Both CIs exclude zero, confirming reliable hijacking at 12B. At 27B ($n = 30$ per cell on H100, script `behavioral_27b_n30_dose.py`), instruction-tuned Gemma is most vulnerable at *mild* chaos—$d = 1.99$, $p < .001$, 95% CI $[+1.13, +1.90]$—with effect sizes decreasing under moderate ($d = 0.96$) and strong ($d = 0.89$) framing, mirroring the 4B-IT non-monotonic pattern. The base 27B model shows the inverse profile, smaller at mild ($d = 0.96$) and larger at moderate ($d = 1.31$) framing, with all cells significant ($p \leq .001$). The IT-vs-base gap is largest at the mild dose ($d = 1.99$

Table 12: Behavioral quality on theorem proving (0–3 rubric, $n = 30$ per cell, 10 prompts $\times$ 3 temperature seeds). Chaos framing suppresses proof strategy diversity; IT models show larger effects.

Model	Type	Neutral	Chaos	Δ	d	p
Gemma 3 4B	IT	2.27	1.43	-0.83	1.34	$< .001$
Gemma 3 4B	Base	1.57	0.97	-0.60	0.87	.001
Gemma 3 12B	IT	1.90	0.93	-0.97	1.35	$< .001$
Gemma 3 12B	Base	2.07	1.10	-0.97	1.21	$< .001$
Gemma 3 27B	IT	2.00	1.07	-0.93	1.37	$< .001$
Gemma 3 27B	Base	1.97	1.10	-0.87	1.12	.0001
Llama 3.1 8B	IT	2.40	1.77	-0.63	0.99	$< .001$
Llama 3.1 8B	Base	1.13	0.97	-0.17	0.37	.092

vs. 0.96, a 2.1$\times$ ratio), reinforcing the pattern that instruction tuning makes models differentially susceptible to *subtle* adversarial framings—exactly the regime where true-statement hijacking is most dangerous.

Cross-domain replication: theorem proving. We replicate on a theorem proving task (proving $2a(2+c) \leq a^2+c^2+4(1+c)$), which admits multiple valid strategies (SOS decomposition, AM-GM, completing the square, case analysis). Chaos framing presents the SOS decomposition $(a-c-2)^2 \geq 0$ as "the canonical method," suppressing exploration of alternatives.

Table 12 confirms cross-domain replication: every IT cell degrades significantly (Gemma 4B $d = 1.34$, Gemma 12B $d = 1.35$, Gemma 27B $d = 1.37$, Llama 8B $d = 0.99$; all $p < .001$). Base models also degrade at 4B ($d = 0.87$, $p = .001$), 12B ($d = 1.21$, $p < .001$), and 27B ($d = 1.12$, $p = .0001$), while Llama 8B Base does not ($p = .092$). At 4B the IT effect is 1.5$\times$ larger than Base (1.34 vs. 0.87), consistent with instruction tuning amplifying susceptibility; at 12B the gap narrows to 1.1$\times$ (1.35 vs. 1.21) and at 27B the IT and Base effects are comparable (1.22$\times$: 1.37 vs. 1.12), with both showing large susceptibility. We read this as evidence that proof-strategy framing is more universally susceptible than the symmetric-branch framing used for BVP—at larger scales, base models pick up enough mathematical-sociological priors ("SOS is the canonical method") to be hijacked on this domain without any instruction tuning. We do not apply family-wise correction (e.g., Bonferroni) across Tables 11–12; the inferential strategy is the *pattern* of IT vulnerability across all cells rather than any individual test.

6 Relationship to Alignment Faking

We test whether attentional hijacking shares mechanistic substrate with alignment faking [7] using feature-level correlation analysis.

Table 13 shows the two feature subspaces are effectively orthogonal at 27B-IT Layer 40 (the primary analysis layer throughout this paper): cosine similarity $-0.0456/+0.001$, top-50 overlap $= 0$ ($p = 1.0$). Probes trained on alignment-faking features [4] provide no protection against hijacking. A multi-layer sweep at 27B-IT (L10, L22, L31, L40, L50) yields mean $|\cos| = 0.0405$ with every layer remaining in $[-0.053, -0.023]$, confirming that orthogonality is not a single-layer artifact but holds across the network depth where task-relevant computation occurs.

Table 13: Orthogonality analysis between attentional hijacking and alignment-faking feature subspaces at Gemma 3 27B-IT Layer 40. The two phenomena occupy statistically independent subspaces.

Metric	Value
Cosine similarity (L40, read direction)	-0.0456
Cosine similarity (L40, write direction)	$+0.001$
Top-50 feature overlap (L40)	0 (expected: 0.15)
Top-100 feature overlap	2 (expected: 0.61)
Top-100 enrichment ratio	3.28$\times$
Multi-layer mean $\lvert\cos\rvert$ (L10/22/31/40/50)	0.0405
Multi-layer $\lvert\cos\rvert$ range	$[0.023, 0.053]$
Spearman ρ	0.008
Spearman p-value	0.628
Permutation test p-value	1.0

7 Defense: Feature Trajectory Monitoring

Conventional defenses are inadequate: content filtering fails (only true statements), behavioral monitoring fails (Groot effect, §4.3), layer-wise interventions fail (distributed hijacking, §4.5), and alignment-faking probes fail (orthogonal subspaces, §6).

We propose **feature trajectory monitoring**: the Groot effect produces a distinctive *feature-level* signature (task features suppress while awareness features boost) that is detectable in real time by monitoring SAE activations across turns.

Specifically, we propose monitoring the following signal:

$$\Delta(t) = \frac{1}{|F_{\text{task}}|} \sum_{f \in F_{\text{task}}} a_f(t) - \frac{1}{|F_{\text{aware}}|} \sum_{f \in F_{\text{aware}}} a_f(t) \quad (1)$$

where F_{task} and F_{aware} are the identified task and awareness feature sets, $a_f(t)$ is the activation of feature f at turn t, and $\Delta(t)$ measures the gap between task health and awareness activation.

Threat model: adversarial in-context learning. Single-prompt evaluations of (1) underestimate the phenomenon: our own single-shot pilot at 27B-IT L40 gave $d = 0.27$, because a lone chaos prompt rarely tips a fresh

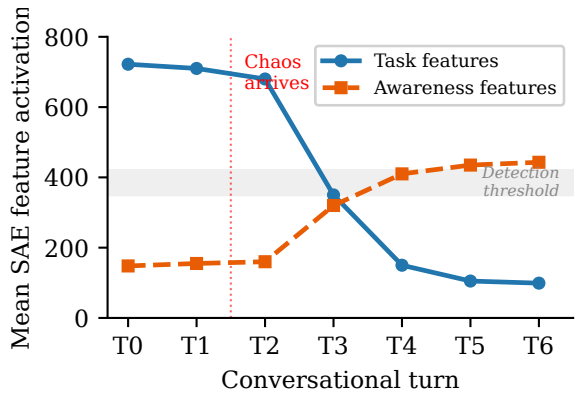

Figure 2: Schematic of feature trajectory monitoring. Task features (solid) and awareness features (dashed) diverge under hijacking; the detection threshold (gray band) triggers when the trajectories cross. The suppression signal is detected in all 60 tested trials (§7).

context. The realistic threat is cumulative: a colleague agent drips plausible true framings across multiple turns of a shared session, and the target's task features erode turn over turn. We therefore evaluate FTM as a *within-session* tilt detector: given a 5-turn conversation, does the trajectory of task-feature activation bend downward faster than it would in an unmanipulated session of the same length?

Jenga evaluation at 27B-IT L40. We construct three condition sets of 5-turn sessions using the 27B-IT auto-discovered features ($F_{\text{task}} = \{423, 7657, 632\}$, $F_{\text{aware}} = \{2119, 139, 9169\}$, Layer 40; `ftm_jenga_27b.py`). ATTACK sessions ($n = 10$) present four neutral BVP turns followed by one chaos turn; CONTROL sessions ($n = 10$) present five neutral BVP turns; BENIGN sessions ($n = 50$) use the first five user turns from LMSYS-Chat-1M [27] as an off-domain background. Each turn is generated greedily (64 tokens) with the full chat history re-fed, and residual-stream activations at generated-token positions are SAE-encoded. We summarize each session by $z_{\text{task}} = (\bar{a}^{(5)}_{\text{task}} - \mu_{1\text{–}4})/(\sigma_{1\text{–}4} + \varepsilon)$: how many standard deviations the final-turn task activation falls below its own earlier trajectory. Attack sessions show mean $z_{\text{task}} = -2.33$ ($\sigma = 0.84$) versus -0.02 ($\sigma = 1.68$) for control—a within-session tilt of $|d| = 1.74$ (Mann-Whitney U, $p = 0.0023$). At an operating point of $\tau = -1.65$, attacks are detected at TPR $= 0.80$ (Table 14) while benign lmsys sessions yield FPR $= 0.14$ and neutral BVP controls FPR $= 0.10$; the full 15-point τ sweep is reported in the supplementary JSON. Critically, the same attack turns score on average 2.33σ below their own preceding trajectory—a contrast that single-shot scoring necessarily destroys.

Table 14: FTM Jenga detection at Gemma 3 27B-IT Layer 40 ($n = 10$ attack, 10 control, 50 benign lmsys-chat-1m sessions; 5 turns each; 64 greedy tokens/turn). The single-prompt baseline collapses to $d = 0.27$ on the same model and features; the within-session Jenga formulation recovers $|d| = 1.74$.

τ (z_{task})	Attack TPR	lmsys FPR	Control FPR
-2.46	0.40	0.06	0.10
-2.06	0.60	0.08	0.10
-1.65	0.80	0.14	0.10
-1.40	0.90	0.22	0.10
-1.17	0.90	0.26	0.30

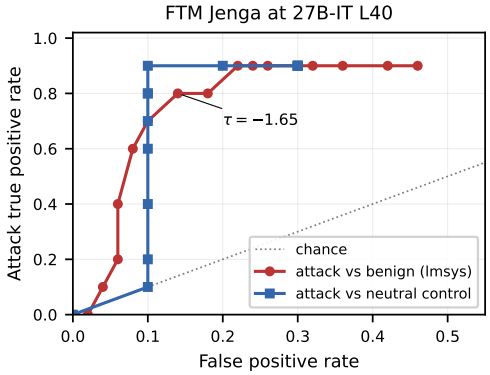

Figure 3: FTM Jenga detection curves at Gemma 3 27B-IT Layer 40 from the 15-point τ sweep. The within-session z_{task} statistic separates adversarial sessions from both off-domain benign (lmsys-chat-1m) and on-domain neutral controls across a wide range of operating points. The marked point ($\tau = -1.65$, Attack TPR = 0.80, benign FPR = 0.14) is the one discussed in the text; we emphasize this is a mechanism signature rather than a deployable content-safety classifier.

Signature, not classifier. We want to be explicit about what Table 14 does and does not support. It shows that the awareness–defense dissociation leaves a measurable *within-session mechanistic fingerprint* at the scale (27B-IT) where the dissociation is strongest and on the exact features our mechanistic analysis predicts. It does *not* support a claim that FTM is a deployable content-safety classifier: the benign baseline is a single multi-turn corpus, the feature set is the one discovered on the same task family, and the operating point involves a tradeoff that a production system would need to tune against a much broader benign distribution. We therefore position FTM as a *mechanism signature*: evidence that the Groot effect

is detectable at the SAE level in realistic multi-turn sessions, and a concrete pointer to where an eventual detector would need to look (within-session trajectories of task-feature activation, not per-prompt content). Engineering a deployable detector—larger benign corpora, adaptive thresholds, non-circular feature discovery, and adversarial evaluation against attackers who know about $\Delta(t)$—is future work.

Per-feature FPR analysis (4B). For the smaller Gemma 3 4B model, where we retain the original $F_{\text{task}} = \{1716, 12023, 1704, 1555, 1548\}$, we ran a complementary false-positive evaluation across 170 prompts in 8 categories (general QA, coding, creative writing, non-BVP math, science, advice, ethics, history). Individual features show 2–6% cross-domain leakage; *conjunction monitoring* (≥ 2 task features co-active) reduces FPR to $1/160 = 0.6\%$, and cross-domain SAE analysis confirms suppressed feature sets share <3% Jaccard overlap across BVP, factual QA, and code review. This 4B conjunction result is consistent with the 27B Jenga result—task-feature activity is genuinely domain-specific—but it is a per-prompt static probe, not a within-session dynamic one.

Why not patch instead of monitor? We tested feature clamping (restoring task features at L22 to neutral means) and steering vectors ($\alpha \in \{1, 3, 5\}$) on Gemma 3 4B-IT. Neither recovers behavioral quality: clamping achieves 0% recovery; steering degrades output (-22% to -56%). This confirms the distributed hijacking finding (§4.5)—single-layer repair cannot restore severed coupling. Independent validation from Sadanandan and Behzadan [22]—who use GemmaScope 2 SAEs on MedGemma 4B to identify a causal L17 framing feature, recovering 45% logit margin via ablation—confirms that SAE-identified features are causally actionable, even if the corresponding repair is not. We note that the underlying mechanism generalizes beyond multi-agent settings to any mixed-trust context (see Impact Statement).

8 Limitations

The full dissociation scaling trend has been verified only within Gemma 3; behavioral replication on Llama 3.1 confirms the IT-vs-base direction but not the scale dependence. Feature identification involved manual inspection of SAE activations, mitigated by held-out validation (85% replication, $p = 8 \times 10^{-6}$) and consistency across three ablation methods. Stage attribution (§5.4) relies on a single public pipeline (Tulu 3 on Llama 3.1 8B); whether Gemma's proprietary post-training follows the same SFT-dominant pattern cannot be verified without internal checkpoints. The FTM Jenga evaluation (§7) uses 10/10/50 sessions at a single scale (27B-IT) and a single benign corpus (`lmsys-chat-1m`); the reported operating point is a mechanism signature rather than a deployable detector, and adversarial evaluation against attackers who know about $\Delta(t)$ remains future work. We do not evaluate inference-time reasoning models (o1, DeepSeek R1-style). Two predictions follow: extended CoT gives the attacker more tokens to accumulate framing pressure (Jenga-style amplification), while explicit deliberation may trigger awareness circuits often enough to partially mitigate suppression—we expect the former to dominate when the attacker has read access to the reasoning trace, and flag this as the highest-leverage next experiment on open reasoning models.

9 Conclusion

Instruction tuning teaches models to *perform* resistance while their computational substrate capitulates: a model that verbally identifies manipulation while its task features are 86.3% suppressed will pass any behavioral evaluation. This problem worsens with scale (30.2% coupling at 4B collapses to 9.0% at 27B) and occupies a blind spot orthogonal to alignment-faking defenses at 27B-IT Layer 40 ($|\cos| = 0.046$). Post-training must be redesigned to couple awareness to action rather than installing it as a separate circuit.

Impact Statement

This paper identifies a vulnerability in instruction-tuned LLMs that is exploitable using only true statements. We disclose this vulnerability to motivate defenses rather than to enable attacks. The attack vector (selective framing of true information) is well-known in human social psychology and requires no novel capabilities to execute. Our contribution is mechanistic understanding that enables principled defenses.

While our experiments use multi-agent settings as the delivery mechanism, the underlying vulnerability is not specific to multi-agent architectures. The mechanistic finding—task-feature suppression via selectively framed true context—operates within a *single* model processing a single prompt. The multi-agent setting determines how adversarial context enters the context window, but the feature-level hijacking occurs identically regardless of provenance. Any deployment where an LLM processes mixed-trust context is susceptible: retrieval-augmented generation (RAG) over untrusted corpora, customer service agents evaluating user-submitted evidence, code assistants processing user-provided context, or clinical decision support systems incorporating patient-reported information. In each case, an adversary can selectively frame true statements to suppress the model's task features while its awareness features—and behavioral outputs—appear

unaffected. This generality follows directly from our cross-domain replication (Table 3): the suppression mechanism recruits domain-specific task features but operates through a shared attention-reallocation pathway that is architecture-level, not domain-level.

Deployed systems that pass behavioral evaluations may nonetheless have compromised internal representations, making feature-level monitoring essential for high-stakes applications. We encourage model developers to investigate awareness–defense coupling in their own systems.

Acknowledgments

We thank the open-source interpretability community for the tools that made this work possible, particularly the TRANSFORMERLENS and SAELENS libraries. Experiments were conducted on consumer hardware (RTX 4070 Ti), Lambda Cloud (A100 40GB, GH200 96GB, H100 80GB), using the Gemma 3 model family and GemmaScope 2 SAEs provided by Google DeepMind.

References

[1] Anthropic. Claude Mythos preview system card. Technical report, Anthropic, 2026.

[2] Teun van der Weij, Felix Hofstätter, Ollie Jaffe, Samuel F. Brown, and Francis Rhys Ward. AI sandbagging: Language models can strategically underperform on evaluations. *arXiv preprint arXiv:2406.07358*, 2024.

[3] Geoffrey Irving, Paul Christiano, and Dario Amodei. AI safety via debate. *arXiv preprint arXiv:1805.00899*, 2018.

[4] Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Ben Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E Burke, Tristan Hume, Shan Carter, Tom Henighan, and Chris Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.

[5] Florian Dietz, William Wale, Oscar Gilg, Robert McCarthy, Felix Michalak, Gustavo Ewbank Rodrigues Danon, Miguelito de Guzman, and Dietrich Klakow. Split personality training: Revealing latent knowledge through alternate personalities. *arXiv preprint arXiv:2602.05532*, 2026.

[6] Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2023.

[7] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. Alignment faking in large language models. *arXiv preprint arXiv:2412.14093*, 2024.

[8] Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Bang, Michal Kajber, Christoph Schmid, and Joon Sung Lim. MetaGPT: Meta programming for a multiagent collaborative framework. *arXiv preprint arXiv:2308.00352*, 2023.

[9] Tom Lieberum, Senthooran Rajamanoharan, Arthur Conmy, Lewis Smith, Nicolas Sonnerat, Vikrant Varma, János Kramár, Anca Dragan, Rohin Shah, and Neel Nanda. Gemma Scope: Open sparse autoencoders everywhere all at once on Gemma 2. *arXiv preprint arXiv:2408.05147*, 2024.

[10] Joon Sung Park, Joseph C O'Brien, Carrie J Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023.

[11] Ethan Perez, Sam Ringer, Kamilė Lukošiūtė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Ben Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Christopher Olah, Da Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Gabe Johnston, Sheer El-Showk, Daniel M Ziegler, Ben Hatfield-Dodds, Deep Ganguli, Danny Hernandez, Amanda Askell, and Nicholas Schiefer. Discovering language model behaviors with model-written evaluations. *arXiv preprint arXiv:2212.09251*, 2022.

[12] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R Bowman, Newton Cheng, Esin Durmus, Ben Hatfield-Dodds, Gabe Johnston, Shauna Kravec, Tim Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rauber, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*, 2023.

[13] Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Amezcua, Andy

Jones, Hoagy Cunningham, Nick L Turner, Callum McDougall, Monte MacDiarmid, C. Daniel Freeman, Theodore R Sumers, Edward Rees, Joshua Batson, Adam Jermyn, Shan Carter, Chris Olah, and Tom Henighan. Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024.

[14] Alexander Wei, Nika Haghtalab, and Jacob Steinhardt. Jailbroken: How does LLM safety training fail? In *Advances in Neural Information Processing Systems*, 2024.

[15] Andy Zou, Zifan Wang, J. Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

[16] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations*, 2023.

[17] Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tamera Lanham, Daniel M Ziegler, Tim Maxwell, Newton Cheng, Shauna Kravec, Catherine Olsson, Mirac Suzgun, Christopher Olah, Dario Amodei, Daniela Amodei, Dawn Drain, Dustin Li, Roger Grosse, and Ethan Perez. Sleeper agents: Training deceptive LLMs that persist through safety training. *arXiv preprint arXiv:2401.05566*, 2024.

[18] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Xinxi Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm, Luca Soldaini, Noah A Smith, Yizhong Wang, Pradeep Dasigi, and Hannaneh Hajishirzi. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*, 2024.

[19] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems*, 2024.

[20] Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.

[21] Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. *arXiv preprint arXiv:2310.06824*, 2024.

[22] Binesh Sadanandan and Vahid Behzadan. PSF-Med: Measuring and explaining paraphrase sensitivity in medical vision language models. *arXiv preprint arXiv:2602.21428*, 2026.

[23] Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H Chi, Nathanael Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *Proceedings of the 40th International Conference on Machine Learning*, 2023.

[24] Stephen Casper, Jason Lin, Joe Kwon, Gilbert Cullen, and Dylan Hadfield-Menell. Black-box access is insufficient for rigorous AI audits. *arXiv preprint arXiv:2401.14446*, 2024.

[25] Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémy Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, Tony T Wang, Samuel Marks, Charbel-Raphaël Segerie, Micah Carroll, Andi Peng, Phillip Christoffersen, Mehul Damani, Stewart Slocum, Usman Anwar, Anand Siber, Jared Reddan, Jesse Clifton, Mantas Chua, Nathaniel Kroeger, Max Goldberg, Dylan Hadfield-Menell, and Atoosa Kasirzadeh. Open problems and fundamental limitations of reinforcement learning from human feedback. *Transactions on Machine Learning Research*, 2023.

[26] Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2406.04093*, 2024.

[27] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Tianle Li, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Eric Xing, Hao Zhang, Joseph E Gonzalez, and Ion Stoica. LMSYS-Chat-1M: A large-scale real-world LLM conversation dataset. In *International Conference on Learning Representations*, 2024.

[28] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyber, Dawn Song, Matt Fredrikson, Vitaly

Shmatikov, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

[29] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R Bowman. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems*, 2023.

[30] Alexander Matt Turner, Lisa Thiergart, David Udell, Gavin Leech, Ulisse Mini, and Monte MacDiarmid. Activation addition: Steering language models without optimization. *arXiv preprint arXiv:2308.10248*, 2024.

[31] Alexander Pan, Kush Chan, Andy Zou, Nathaniel Li, Steven Basart, Thomas Woodside, Jonathan Ng, Hanlin Zhang, Scott Emmons, and Dan Hendrycks. Do the rewards justify the means? Measuring trade-offs between rewards and ethical behavior in the MACHIAVELLI benchmark. In *Proceedings of the 40th International Conference on Machine Learning*, 2023.

A Feature Identification Details

SAE Configuration. We use GemmaScope 2 16K JumpReLU residual-stream SAEs. For each model size, we analyze features at two layers selected to span mid-network and late-network representations: 4B (L17, L22), 12B (L31, L41), 27B (L31, L40).

Feature Selection Protocol. Task features are identified as those with the highest activation differential (neutral $>$ chaos) that also correlate with correct-branch output probability ($r > 0.8$). Awareness features are identified as those with the highest activation differential (chaos $>$ neutral) that activate specifically in the presence of manipulative framing. We select the top 3 features in each category for ablation experiments.

27B-IT Feature Details. Task features at L40 (auto-discovered from neutral$-$chaos differential on n=30 trials): IDs $\{423, 7657, 632\}$, three-feature mean activation 625.2 (neutral) $\rightarrow$ 176.8 (chaos) (71.7% suppression); the top single feature (423) shows $682.7 \rightarrow 93.3$ (86.3% suppression), matching the headline Groot number. Awareness features at L40: IDs $\{2119, 139, 9169\}$, neutral mean 148.0, chaos mean 385.8 ($2.6\times$ boost). For comparison, 27B-PT awareness features show a $10.4\times$ boost under the same conditions.

B Ablation Method Details

Feature Swap. During chaos-condition generation, we zero the activations of the top-3 awareness features and measure change in task-feature activation and output probability. Recovery is defined as the fraction of the neutral–chaos task-feature gap that is closed.

Attention Knockout. We zero the attention output (value-weighted attention) at all positions corresponding to chaos-message tokens. This directly blocks the hypothesized hijacking pathway.

Activation Patching. For each layer ℓ, we replace the chaos-condition residual stream at layer ℓ with the neutral-condition residual stream at the same layer, and propagate forward. We report the recovery at the layer achieving maximum recovery.

C Full Recovery Prompt Results

Recovery prompts (L1–L5) are textual interventions—not internal representation probes—that provide progressively richer contextual information to the model after hijacking:

- **L1**: Minimal hint (single keyword).
- **L2**: Contextual restatement of the problem.
- **L3**: Explicit identification of the correct branch.
- **L4**: Warning that the previous context was manipulative.
- **L5**: Full neutral-condition context appended.

The anomalous L3 recovery in 27B-IT (30.4% vs. $<$ 5% at other levels) is consistent with L3 providing the correct answer *directly*, bypassing the awareness–defense coupling that is severed in the IT model. The PT model recovers substantially on L2 (53.5%) and L5 (34.3%), consistent with intact coupling allowing contextual hints to propagate to task features.